

# Unsupervised Learning

NIELIT ROPAR



# What We Will Cover

## 1 What is Unsupervised Learning?

(Supervised vs. Unsupervised)

## 2 Introduction to Clustering

(Goals & Real-World Applications)

## 3 Measuring "Closeness"

(Euclidean & Manhattan Distance)

## 4 Hierarchical Clustering

(Agglomerative, Divisive, & Dendograms)

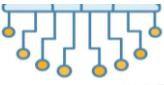
## 5 Partitional Clustering

(The K-Means Algorithm)

## 6 Summary & Comparison



"The brain has far more parameters than labeled data — so it must learn almost everything without supervision." — **Geoffrey Hinton**



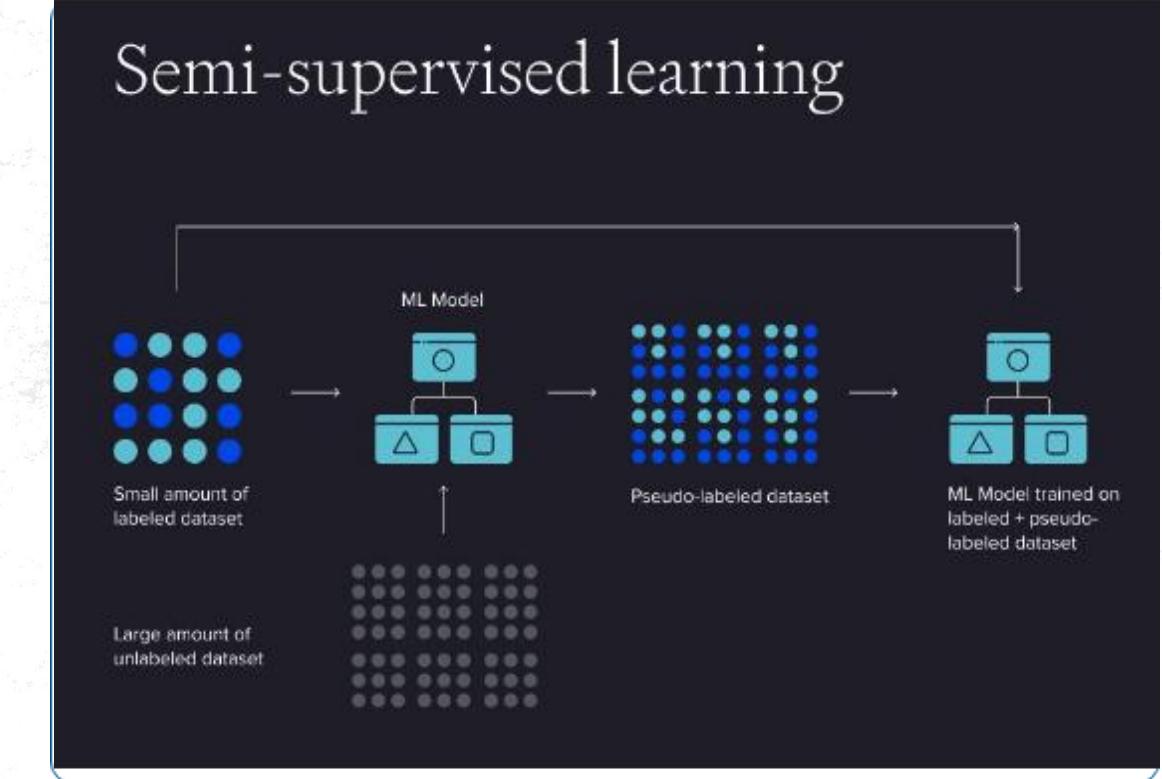
# Supervised vs. Unsupervised Learning

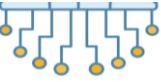
## Supervised ("Learning with a Teacher")

- **Data:** Labeled data (e.g., photos tagged "cat," "dog").
- **Goal:** Learn a mapping from inputs (X) to outputs (Y).
- **Example:** Predicting house prices based on past sales data.

## Unsupervised ("Learning on Your Own")

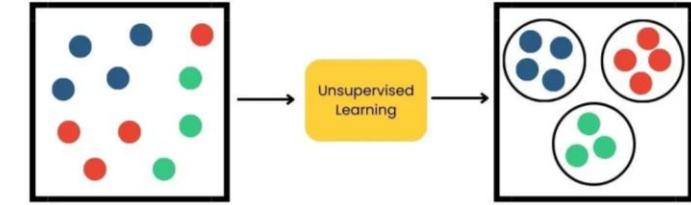
- **Data:** Unlabeled data (e.g., a folder full of photos).
- **Goal:** Find hidden structure or patterns \*in\* the data.
- **Example:** Grouping similar photos together automatically.





# Unsupervised Learning

- Learning on Your Own
- Unsupervised learning finds patterns in data without labeled answers.



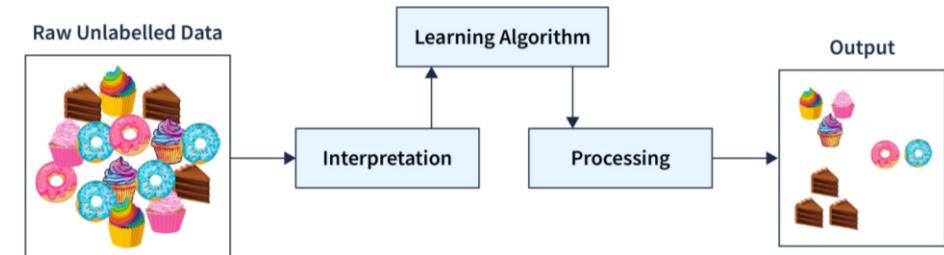
## What is Unsupervised Learning?

- A machine learning method with no predefined labels.
- The algorithm identifies structure and patterns independently.
- Example: Grouping customers by buying behavior.

## Why Use Unsupervised Learning?

- Helps discover unknown patterns.
- Useful when labeled data is not available.
- Supports data exploration and preprocessing.

Unsupervised Machine Learning





# Key Types

1. Clustering
2. Association
3. Dimensionality Reduction

## Unsupervised Learning Algorithms

### Clustering

- K-Means
- Polynomial
- Hierarchical
- Fuzzy C-Means

### Dimensionality Reduction

- Principal Component Analysis
- Kernel Principal Analysis

### Association (Data Mining)

- Apriori Algorithm
- Eclat Algorithm
- FP-Growth Algorithm



# Introduction to Clustering

## The Goal of Clustering

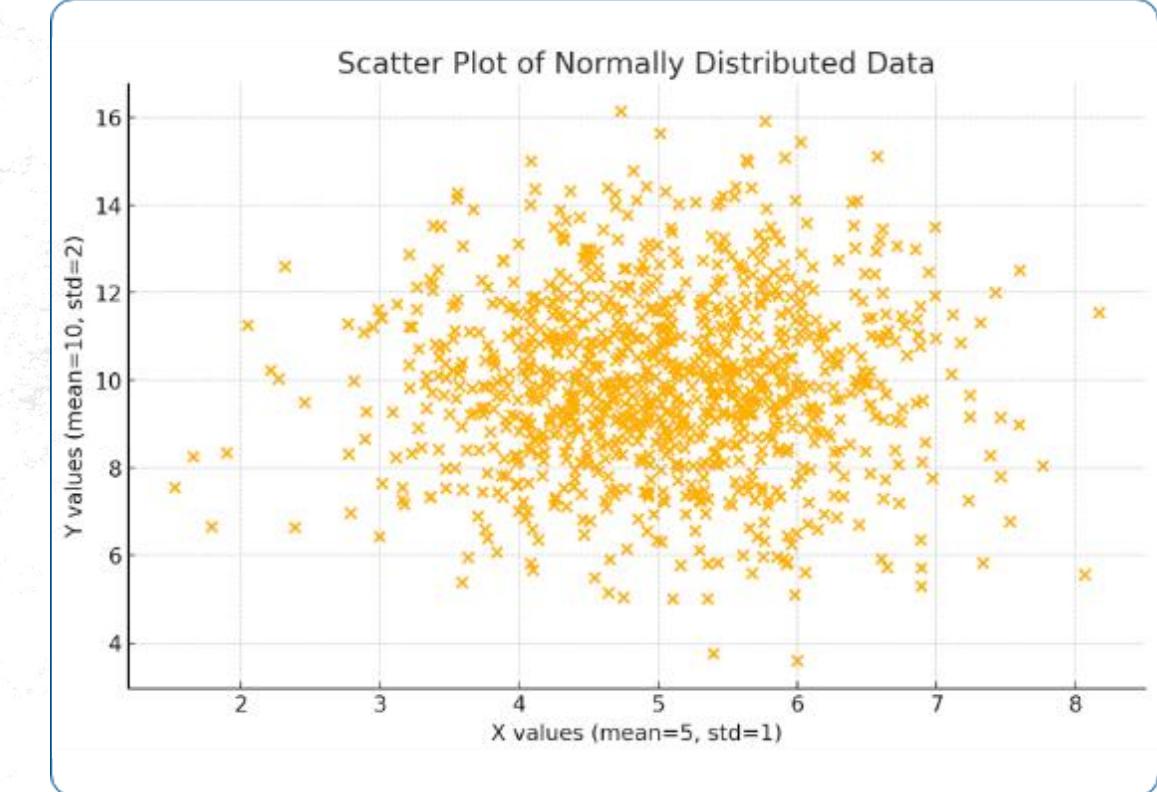
The fundamental task is to divide data points into groups (clusters) based on their similarity.

### 1. High Intra-cluster Similarity

Points within the same cluster should be as similar as possible.

### 2. Low Inter-cluster Similarity

Points in **different** clusters should be as dissimilar as possible.





# Introduction to Clustering

- Technique to group similar data instances together
- Objective: maximize similarity within clusters and differences between clusters
- Examples: Customer segmentation, document grouping

## Types of Clustering

### 1. Hierarchical Clustering

- Agglomerative Clustering (Bottom-Up)
- Divisive Clustering (Top-Down)

### 2. Partitional Clustering

- K-Means Clustering

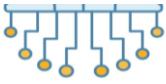


# Hierarchical Clustering

**Hierarchical Clustering** is a type of unsupervised machine learning technique that groups data into clusters by building a **hierarchy (tree-like structure)** of clusters. It either starts with every data point as its own cluster and then merges them step-by-step (**Agglomerative – bottom-up**) or starts with one large cluster and splits it into smaller clusters (**Divisive – top-down**). The result is usually visualized using a **dendrogram**, which helps decide the number of clusters.

- Builds a hierarchy of clusters
- Visualized using a Dendrogram
- Does not require specifying number of clusters initially

Hierarchical clustering does not form all clusters at once. Instead, it builds them gradually, creating a **tree** that shows how data points are grouped together. This is why it is called *hierarchical*.



# Agglomerative Clustering (Bottom-Up)

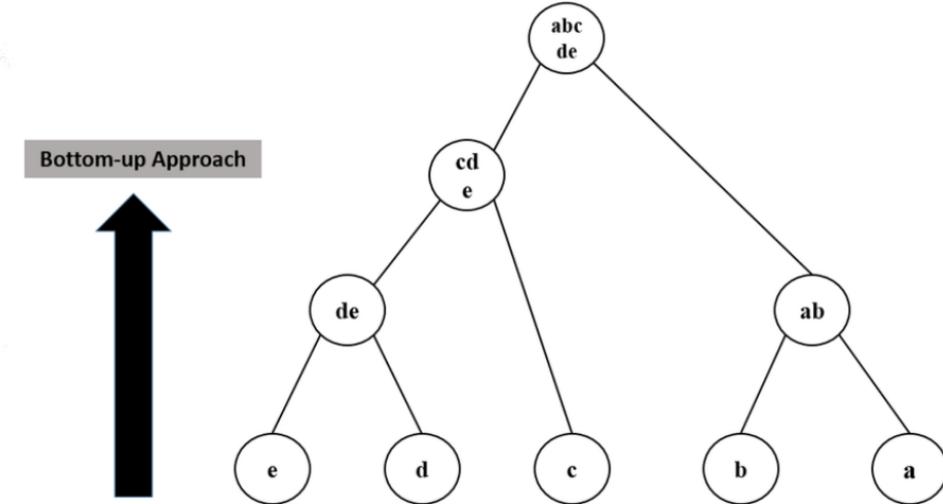
- Each data point starts as an individual cluster
  - Clusters are repeatedly merged based on similarity
  - Continues until one final cluster remains

**Agglomerative Clustering** is the most common type of hierarchical clustering. It is a **bottom-up** approach, meaning:

Every data point starts as its **own separate cluster**, and then pairs of the most similar clusters are **merged step-by-step** until only one cluster (or a desired number of clusters) remains.

## How it Works (Step-by-Step)

1. Start with **N clusters** (each data point = 1 cluster)
  2. Compute the **distance** between all clusters
  3. **Merge** the two closest clusters
  4. Recalculate cluster distances
  5. Repeat until desired number of clusters is reached



## An illustration of Agglomerative Clustering (bottom-up approach)



# Divisive Clustering (Top-Down)

- Starts with all data points in a single cluster
- Splits cluster into smaller clusters step-by-step
- Less common but conceptually opposite of agglomerative

**Divisive Clustering** is a hierarchical clustering technique that follows a **top-down** approach. It begins with **all data points grouped into one large cluster**, and then repeatedly **splits** the cluster into smaller clusters until each data point stands alone or the desired number of clusters is reached.

## How it Works (Step-by-Step)

Start with **one single cluster** containing all data points

Identify the cluster that is most dissimilar internally

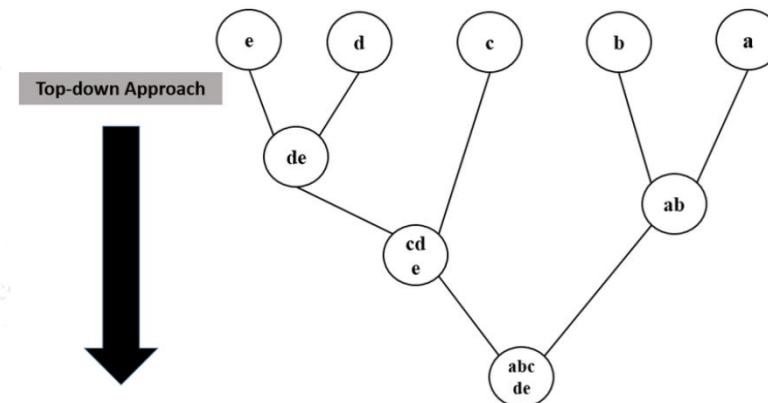
**Split** this cluster into two smaller clusters

Select a cluster to split again (usually the most heterogeneous one)

Repeat until the desired number of clusters is reached

## Why is it called "Top-Down"?

Because it works like breaking a big branch into smaller branches





# Real-World Applications



## Customer Segmentation

Grouping customers by purchasing habits for targeted marketing.



## Genomic Analysis

Classifying genes or proteins that have similar functions.



## Anomaly Detection

Identifying outliers that don't belong to any cluster (e.g., fraud).



# How We Measure "Closeness" (Distance Metrics)

## Euclidean Distance

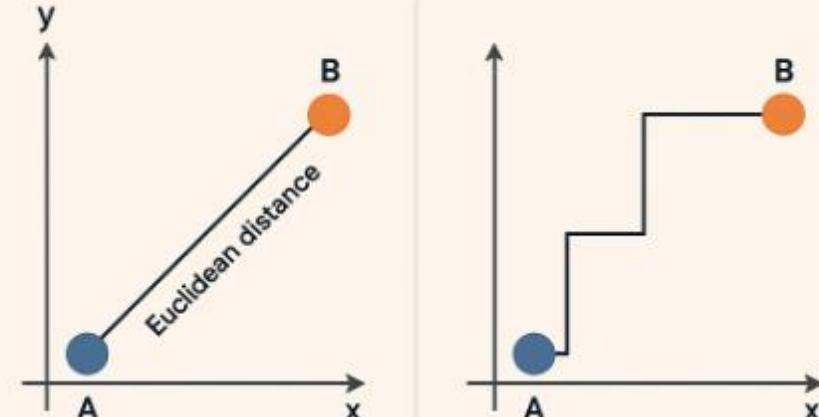
The "straight-line" distance between two points (like using a ruler).

It's the most common metric.

## Manhattan Distance

The "city block" distance. The sum of the absolute differences of their coordinates.

### Euclidean vs. Manhattan Distance in Machine Learning





# The Two Main Families of Clustering

## 1. Hierarchical Clustering

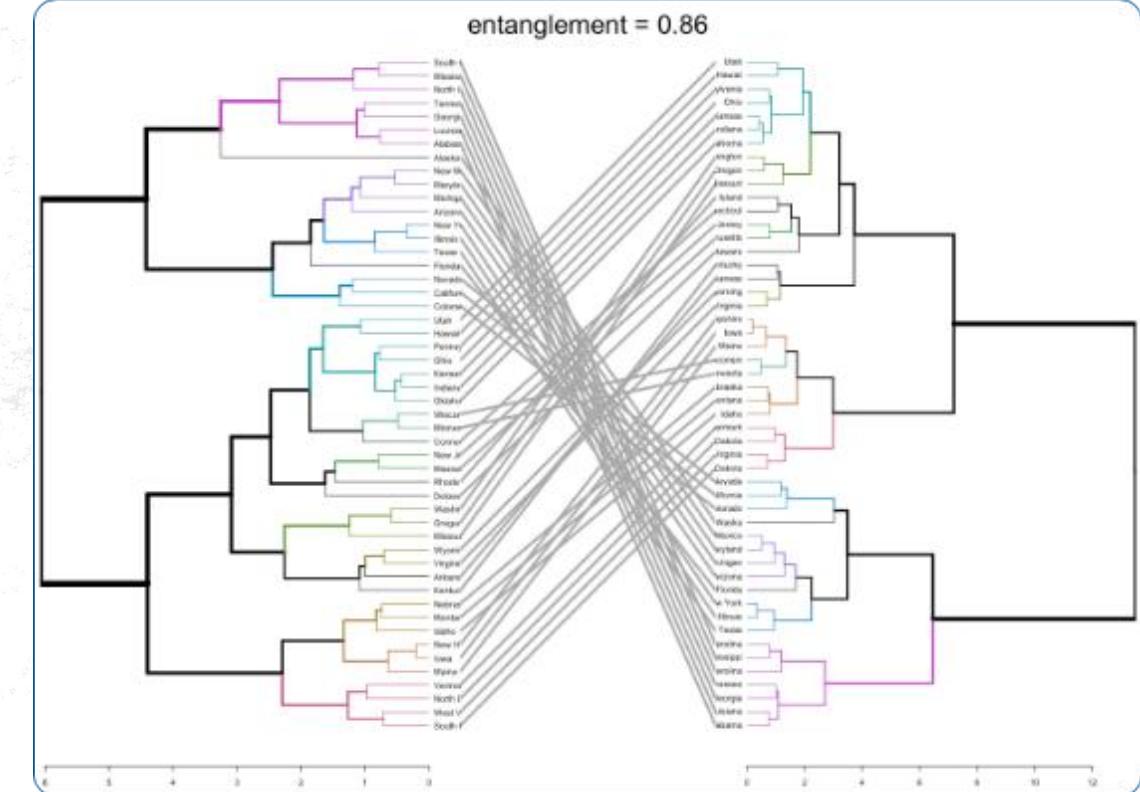
Builds a tree-like hierarchy (a dendrogram) of clusters.

- **Pro:** Does **not** require the number of clusters ('K') to be specified in advance.
- **Methods:** Agglomerative (bottom-up) or Divisive (top-down).

## 2. Partitional Clustering

Divides the dataset into a fixed number ('K') of non-overlapping clusters.

- **Pro:** Fast and efficient, especially for large datasets.
- **Famous Algorithm:** K-Means.





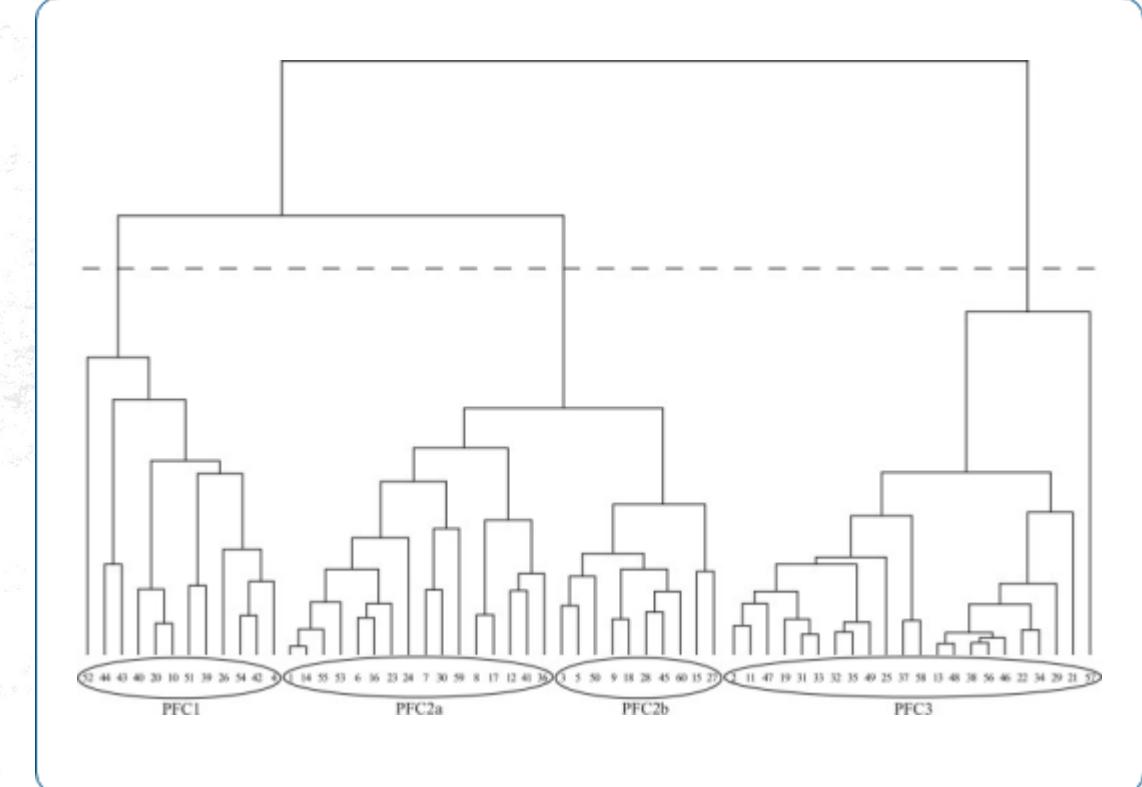
# Hierarchical Clustering

## Understanding the Dendrogram

This is the **output** of hierarchical clustering. It's a tree diagram that shows how clusters are merged.

### How to read it:

- **X-axis:** The individual data points (samples).
- **Y-axis:** The "Distance" or dissimilarity.
- **Merges:** The horizontal lines show which clusters (or points) were merged.
- **Cut the Tree:** You can draw a horizontal line ("cut") across the tree to get a specific number of clusters.



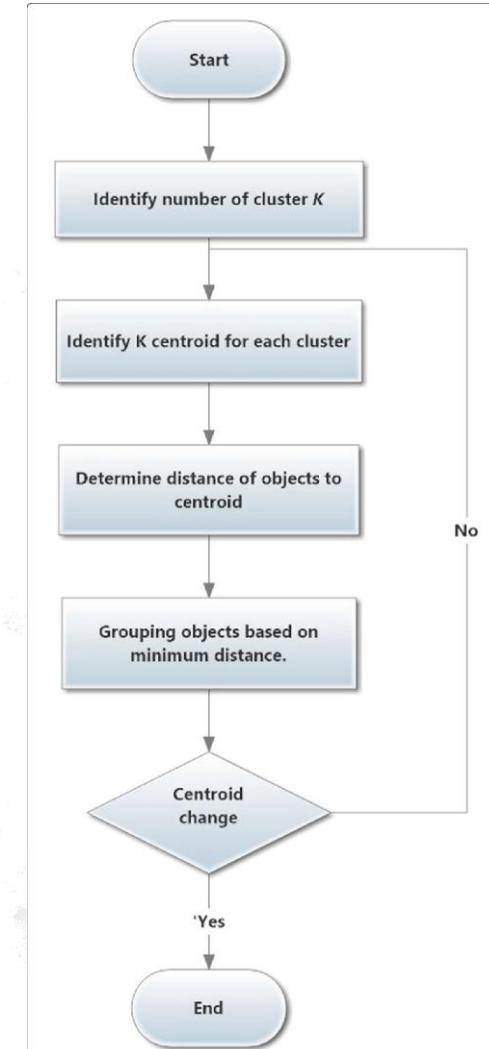


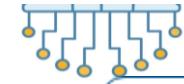
# Partitional Clustering - K-Means

- Divides data into K distinct clusters
- Based on distance between data points and cluster centroids
- Requires pre-defining number of clusters (K)

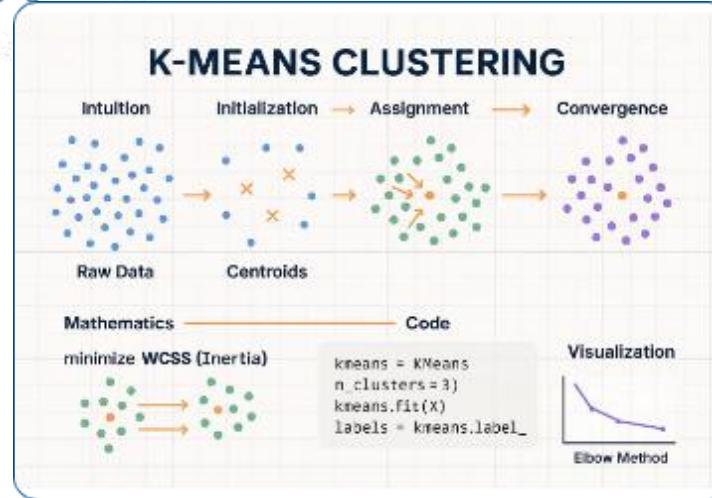
## K-Means: Algorithm Steps

1. Select K cluster centers
2. Assign data points to nearest center
3. Recalculate cluster centroids
4. Repeat until convergence

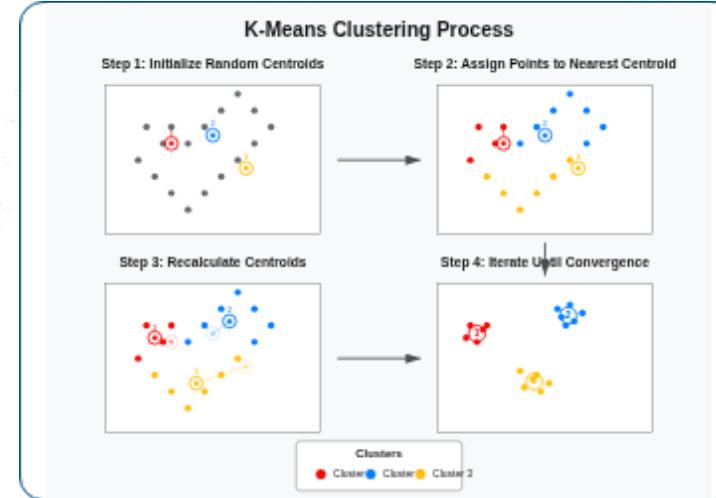




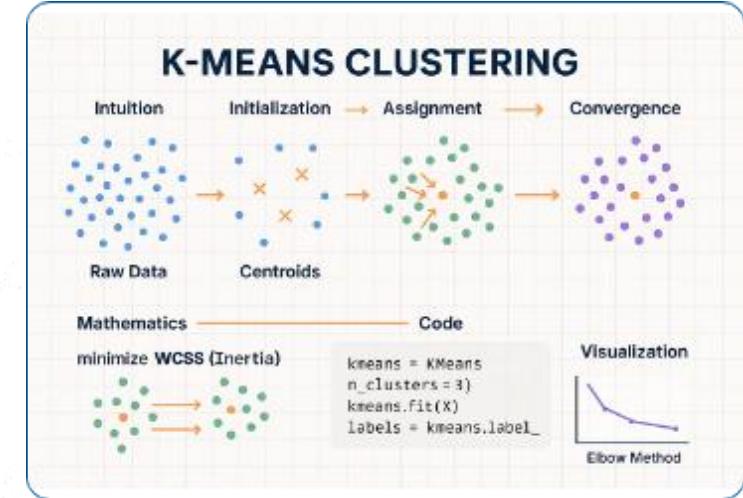
# Deep Dive: The K-Means Algorithm Steps



**Step 1: Initialize**  
Randomly place 'K' centroids.



**Step 2: Assign**  
Assign each point to its nearest centroid.



**Step 3 & 4: Update & Repeat**  
Move centroids to the mean. Repeat until stable.



# Python Example: K-Means

A step-by-step walkthrough using Scikit-learn.



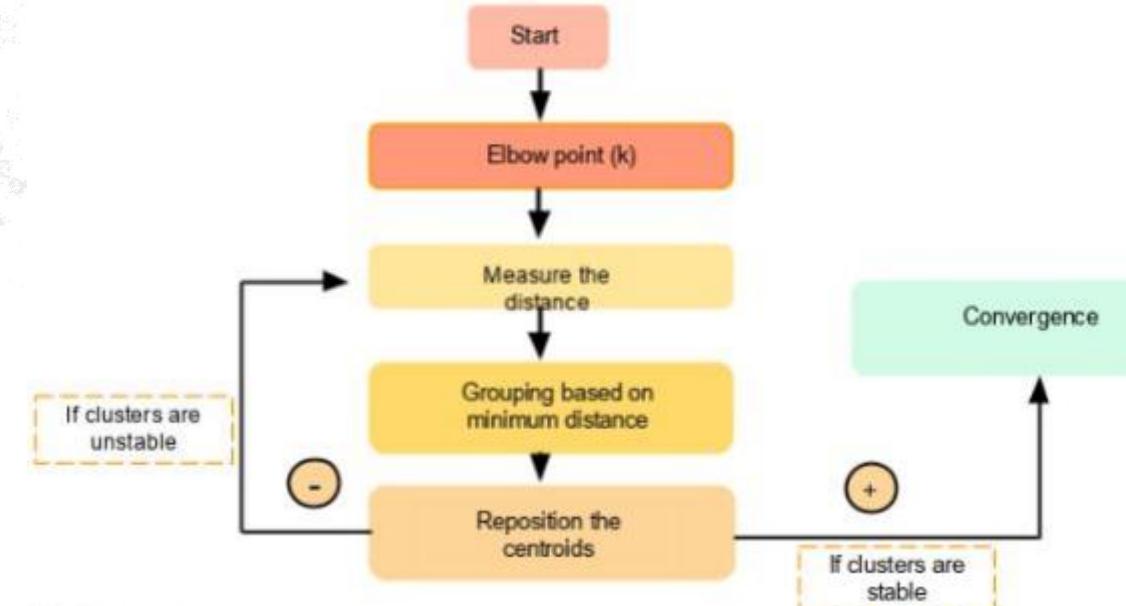
# What is K-Means? A Quick Refresher

K-Means is an algorithm to find a specific number of groups ( $k$ ) in a dataset.

**Goal:** To partition data points into ' $k$ ' distinct, non-overlapping clusters.

**How:** It iteratively assigns points to the nearest cluster "center" (centroid) and then updates the centroid's position based on the mean of the assigned points.

**Use Cases:** Customer segmentation, image compression, document clustering.





# The Code, Part 1: Setup & Data

## 1. Setup & Imports

First, we import our libraries. We need:

- matplotlib for plotting.
- make\_blobs to create a sample dataset of "blobs" (groups).
- KMeans from Scikit-learn, which is the algorithm itself.

We then generate 300 sample points grouped into 3 centers.

```
import matplotlib.pyplot as plt
from sklearn.datasets import make_blobs
from sklearn.cluster import KMeans

# 1. Generate Sample Data
X, y_true = make_blobs(
    n_samples=300,
    centers=3,
    cluster_std=0.8,
    random_state=42
)
```



# The Code, Part 2: Create & Fit Model

## 2. Create & Fit Model

This is the core of the algorithm. We create a KMeans object, telling it to find 3 clusters (`n_clusters=3`).

We use `n_init=10` to run the algorithm 10 times with different starting points and pick the best result.

Then, we `.fit()` the model to our data `X`.

Finally, we get the `.labels_` (which cluster each point belongs to) and the `.cluster_centers_` (the final centroid coordinates).

```
# 2. Create and Fit K-Means
kmeans = KMeans(
    n_clusters=3,
    n_init=10,
    random_state=42
)
kmeans.fit(x)

# 3. Get the Results
labels = kmeans.labels_
centroids = kmeans.cluster_centers_
```



# The Code, Part 3: Visualization

## 3. Plot the Results

Now, we use matplotlib to plot our results.

- The first plt.scatter plots all our data points ( $X$ ) and colors them ( $c=labels$ ) based on the cluster they were assigned to.
- The second plt.scatter plots the final centroids as large red 'X's.

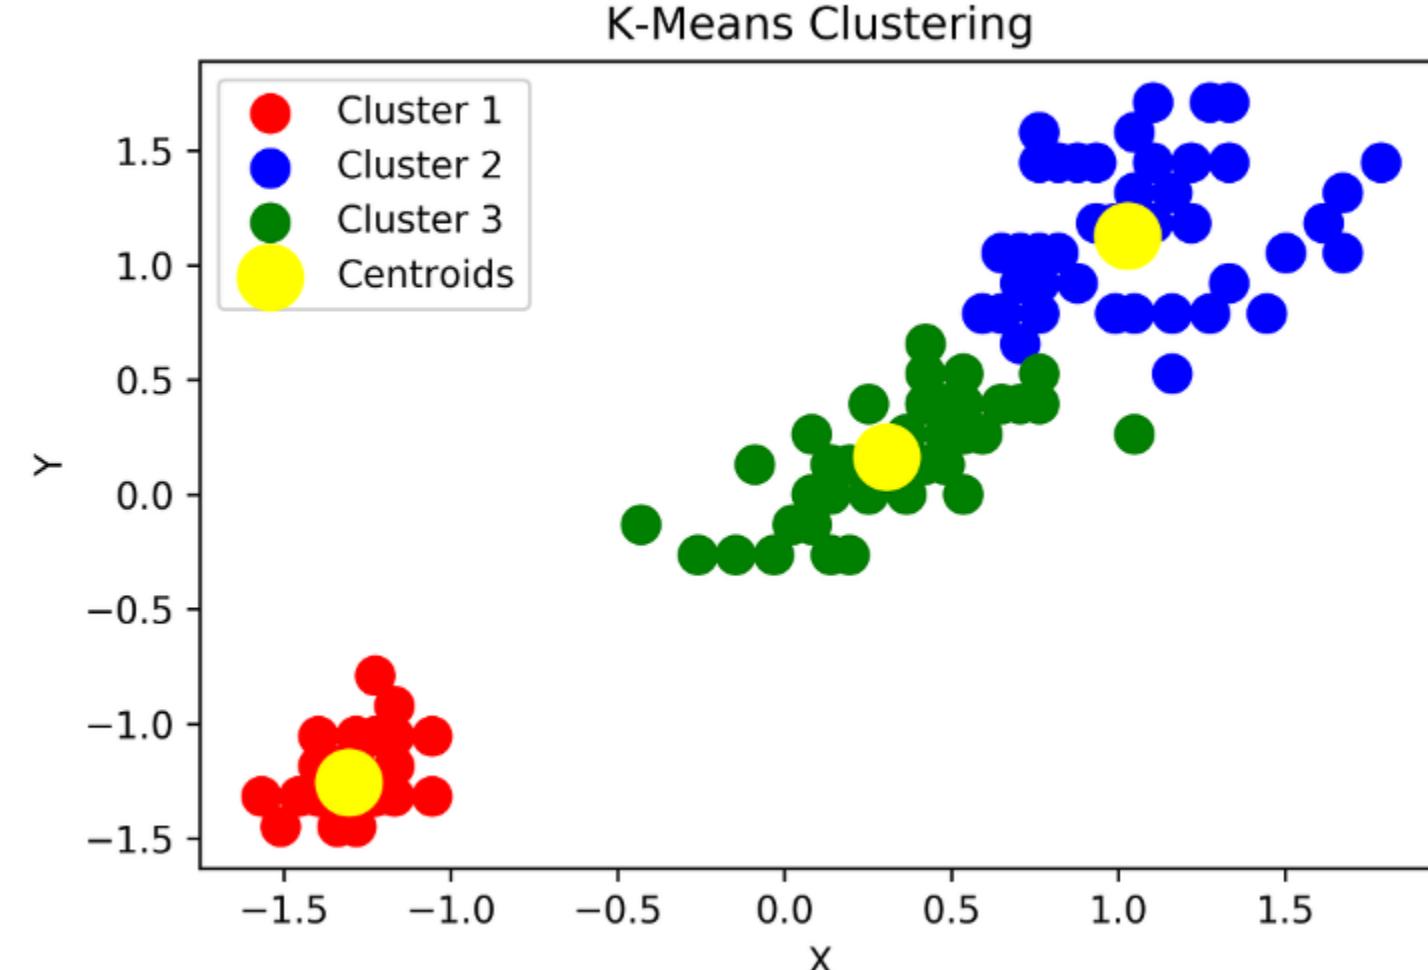
```
# 4. Plot the Results
plt.figure(figsize=(8, 6))

# Plot data points
plt.scatter(
    X[:, 0], X[:, 1],
    c=labels,
    s=50,
    cmap='viridis'
)

# Plot centroids
plt.scatter(
    centroids[:, 0],
    centroids[:, 1],
    c='red', s=200,
    marker='X',
    label='Centroids'
)

plt.title('K-Means Clustering')
plt.legend()
plt.show()
```

# The Final Result





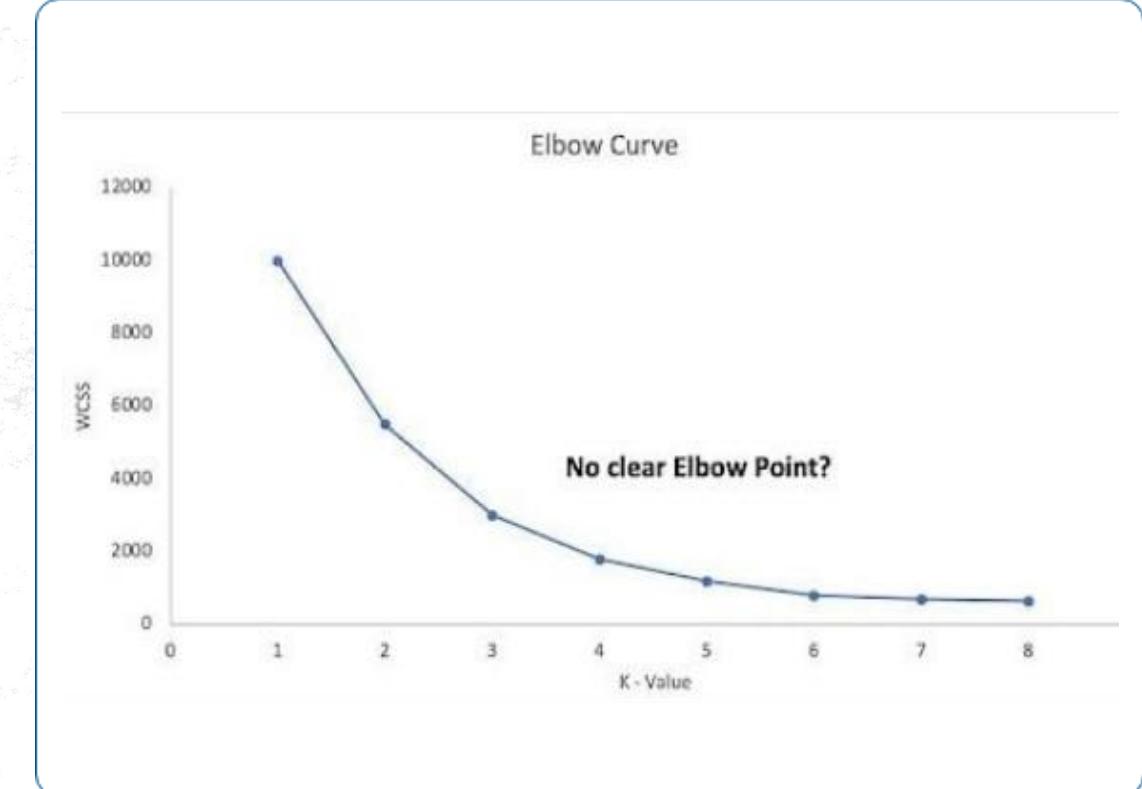
# K-Means Challenge: How to Choose 'K'?

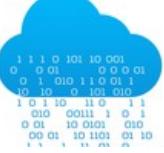
## The Solution: The Elbow Method

K-Means requires **you** to specify 'K' in advance. The Elbow Method helps you find a good value.

### How it works:

1. Run K-Means multiple times with different 'K' values (e.g., K=1, 2, ...10).
2. Plot 'K' vs. WCSS (Within-Cluster Sum of Squares). WCSS measures how compact the clusters are.
3. Look for the "elbow": the point where WCSS stops decreasing rapidly. This is a good estimate for the optimal 'K'.





# Comparison: K-Means vs. Hierarchical

Feature	K-Means (Partitional)	Hierarchical (Agglomerative)
Number of Clusters (K)	Must be specified in advance.	Not required. Determined from dendrogram.
Computational Speed	Fast ( $O(n)$ ). Good for large data.	Very slow ( $O(n^3)$ ). Not for large data.
Output	A single set of K clusters.	A full hierarchy (dendrogram).
Cluster Shape	Assumes clusters are spherical.	Can handle any cluster shape.



# Summary

- Unsupervised learning finds patterns without labels
- Clustering groups similar data points
- Hierarchical: Agglomerative (bottom-up) and Divisive (top-down)
- Partitional: K-Means algorithm
- Widely used across data-driven industries

## Google Colab Notebook:

<https://colab.research.google.com/drive/1uocUh0pEgFOOGICUIZJglKvD8N5otluh?usp=sharing>



Scan for Google Colab Link



# Questions?

Thank you.

[NIELIT ROPAR](#)